

Database System

03 | Indexing

Tahun Ajar Ganjil 2024/2025

Oleh:
Tim Dosen

GOALS OF MEETING

1

Students knows the basic concepts of indexing

2


Students understand how to organize B+-Tree Index Files

3

Students can create indexes on the DBMS to speed up searching performance

OUTLINES

- Basic Concepts
- Ordered Indices
- B+-Tree Index Files
- Creation of Indices



BASIC CONCEPTS



BASIC CONCEPTS

- Indexing mechanisms used to speed up access to desired data.
 - E.g., author catalog in library
- **Search Key** - attribute to set of attributes used to look up records in a file.
- An **index file** consists of records (called **index entries**) of the form:

search-key	pointer
------------	---------

- Index files are typically much smaller than the original file
- Two basic kinds of indices:
 - **Ordered indices:** search keys are stored in sorted order
 - **Hash indices:** search keys are distributed uniformly across “buckets” using a “hash function”.



ORDERED INDICES

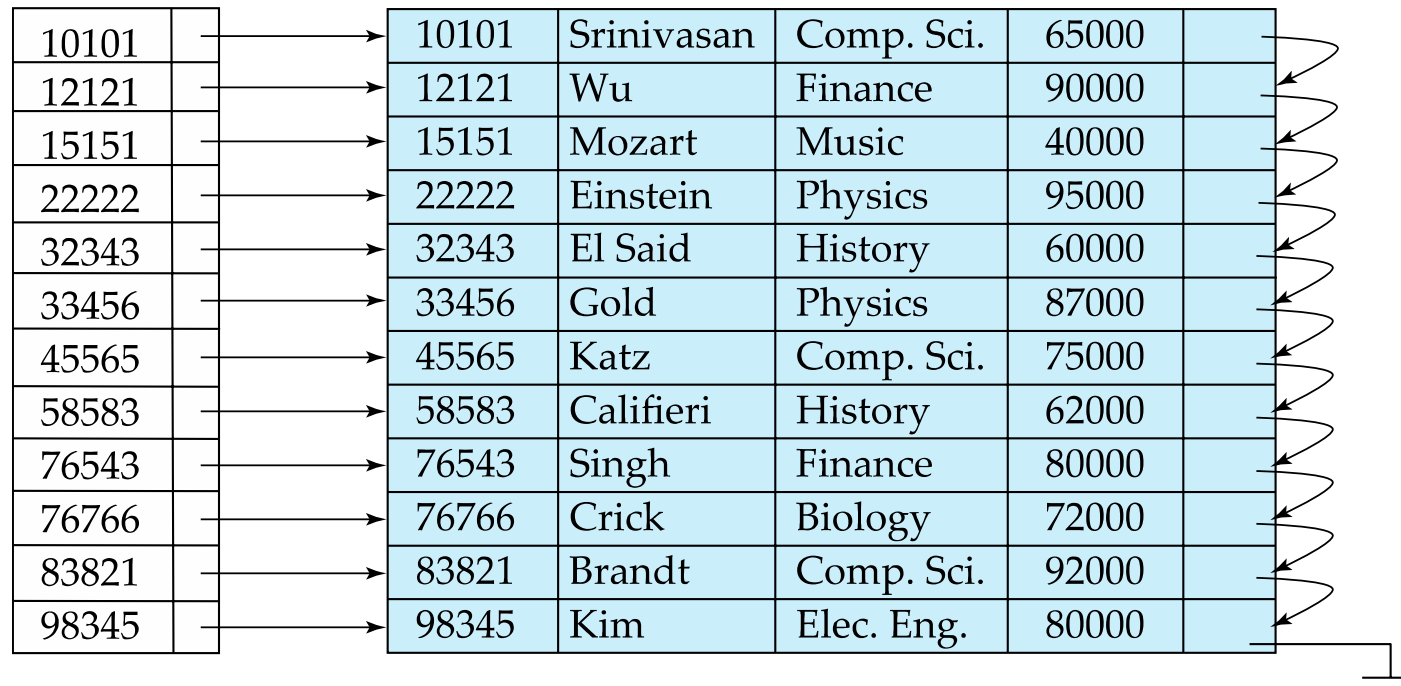
- In an **ordered index**, index entries are stored sorted on the search key value.
- **Clustering index**: in a sequentially ordered file, the index whose search key specifies the sequential order of the file.
 - Also called **primary index**
 - The search key of a primary index is usually but not necessarily the primary key.
- **Secondary index**: an index whose search key specifies an order different from the sequential order of the file. Also called **nonclustering index**.
- **Index-sequential file**: sequential file ordered on a search key, with a clustering index on the search key.



DENSE INDEX FILES

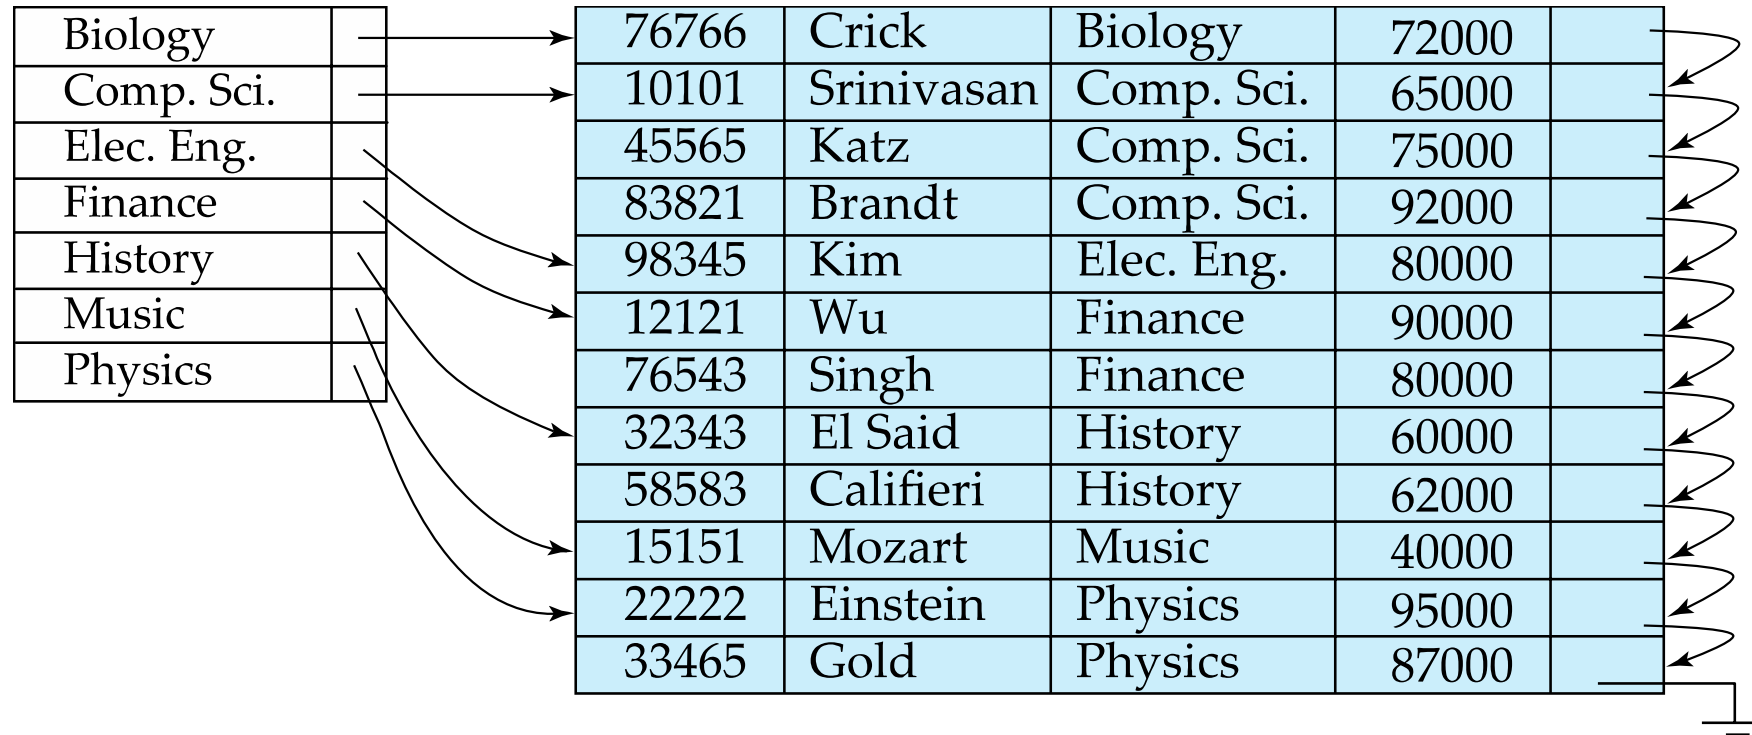
- **Dense index** — Index record appears for every search-key value in the file.
- E.g. index on *ID* attribute of *instructor* relation

10101	→	10101	Srinivasan	Comp. Sci.	65000	→
12121	→	12121	Wu	Finance	90000	→
15151	→	15151	Mozart	Music	40000	→
22222	→	22222	Einstein	Physics	95000	→
32343	→	32343	El Said	History	60000	→
33456	→	33456	Gold	Physics	87000	→
45565	→	45565	Katz	Comp. Sci.	75000	→
58583	→	58583	Califieri	History	62000	→
76543	→	76543	Singh	Finance	80000	→
76766	→	76766	Crick	Biology	72000	→
83821	→	83821	Brandt	Comp. Sci.	92000	→
98345	→	98345	Kim	Elec. Eng.	80000	→



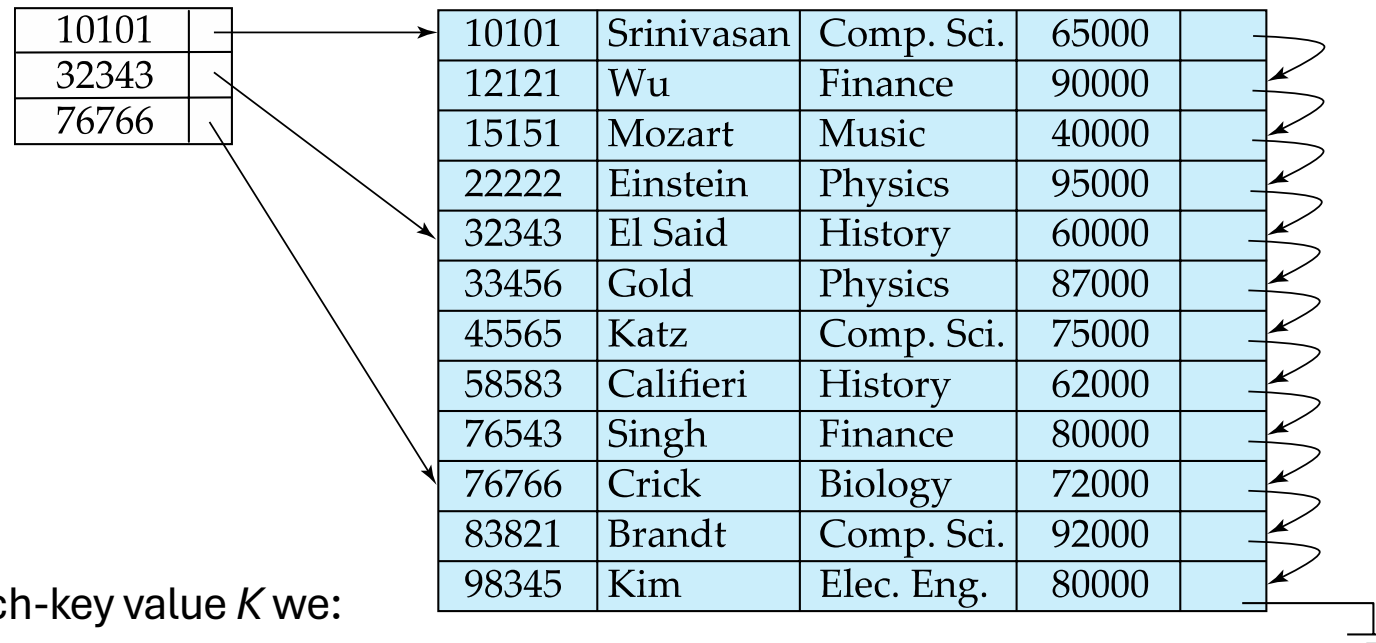
DENSE INDEX FILES (CONT.)

- Dense index on *dept_name*, with *instructor* file sorted on *dept_name*



SPARSE INDEX FILES

- **Sparse Index:** contains index records for only some search-key values.
 - Applicable when records are sequentially ordered on search-key



- To locate a record with search-key value K we:
 - Find index record with largest search-key value $< K$
 - Search file sequentially starting at the record to which the index record points




SPARSE INDEX FILES (CONT.)

- Compared to dense indices:
 - Less space and less maintenance overhead for insertions and deletions.
 - Generally slower than dense index for locating records.

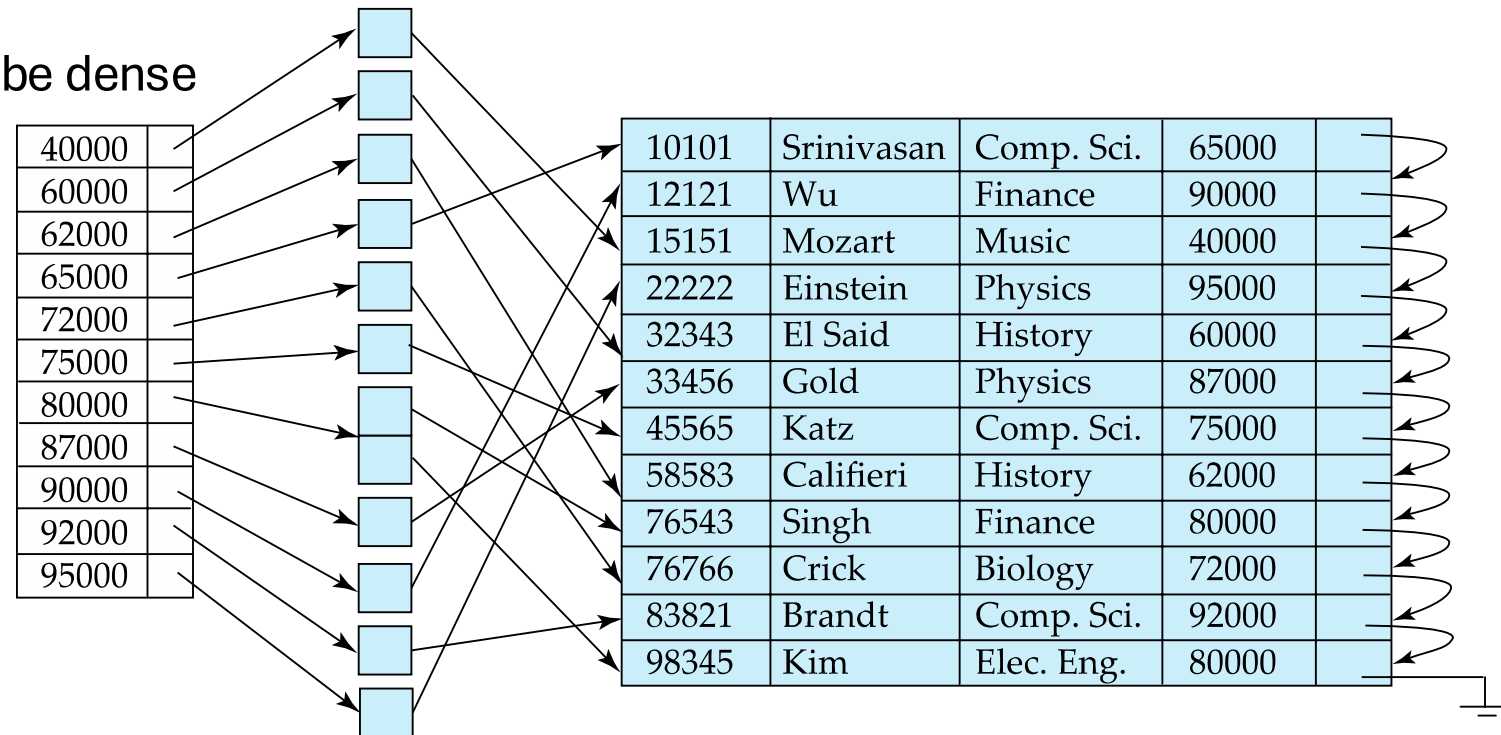
- **Good tradeoff:**

- for clustered index: sparse index with an index entry for every block in file, corresponding to least search-key value in the block.
- For unclustered index: sparse index on top of dense index (multilevel index)



SECONDARY INDICES EXAMPLE

- Secondary index on salary field of instructor
- Index record points to a bucket that contains pointers to all the actual records with that particular search-key value.
- Secondary indices have to be dense



CLUSTERING VS NONCLUSTERING INDICES

- Indices offer substantial benefits when searching for records.
- BUT: indices imposes overhead on database modification
 - when a record is inserted or deleted, every index on the relation must be updated
 - When a record is updated, any index on an updated attribute must be updated
- Sequential scan using clustering index is efficient, but a sequential scan using a secondary (nonclustering) index is expensive on magnetic disk
 - Each record access may fetch a new block from disk
 - Each block fetch on magnetic disk requires about 5 to 10 milliseconds




MULTILEVEL INDEX

- If index does not fit in memory, access becomes expensive.
- Solution: treat index kept on disk as a sequential file and construct a sparse index on it.
 - outer index – a sparse index of the basic index
 - inner index – the basic index file
- If even outer index is too large to fit in main memory, yet another level of index can be created, and so on.
- Indices at all levels must be updated on insertion or deletion from the file.



MULTILEVEL INDEX (CONT.)



B+ - TREE INDEX FILES

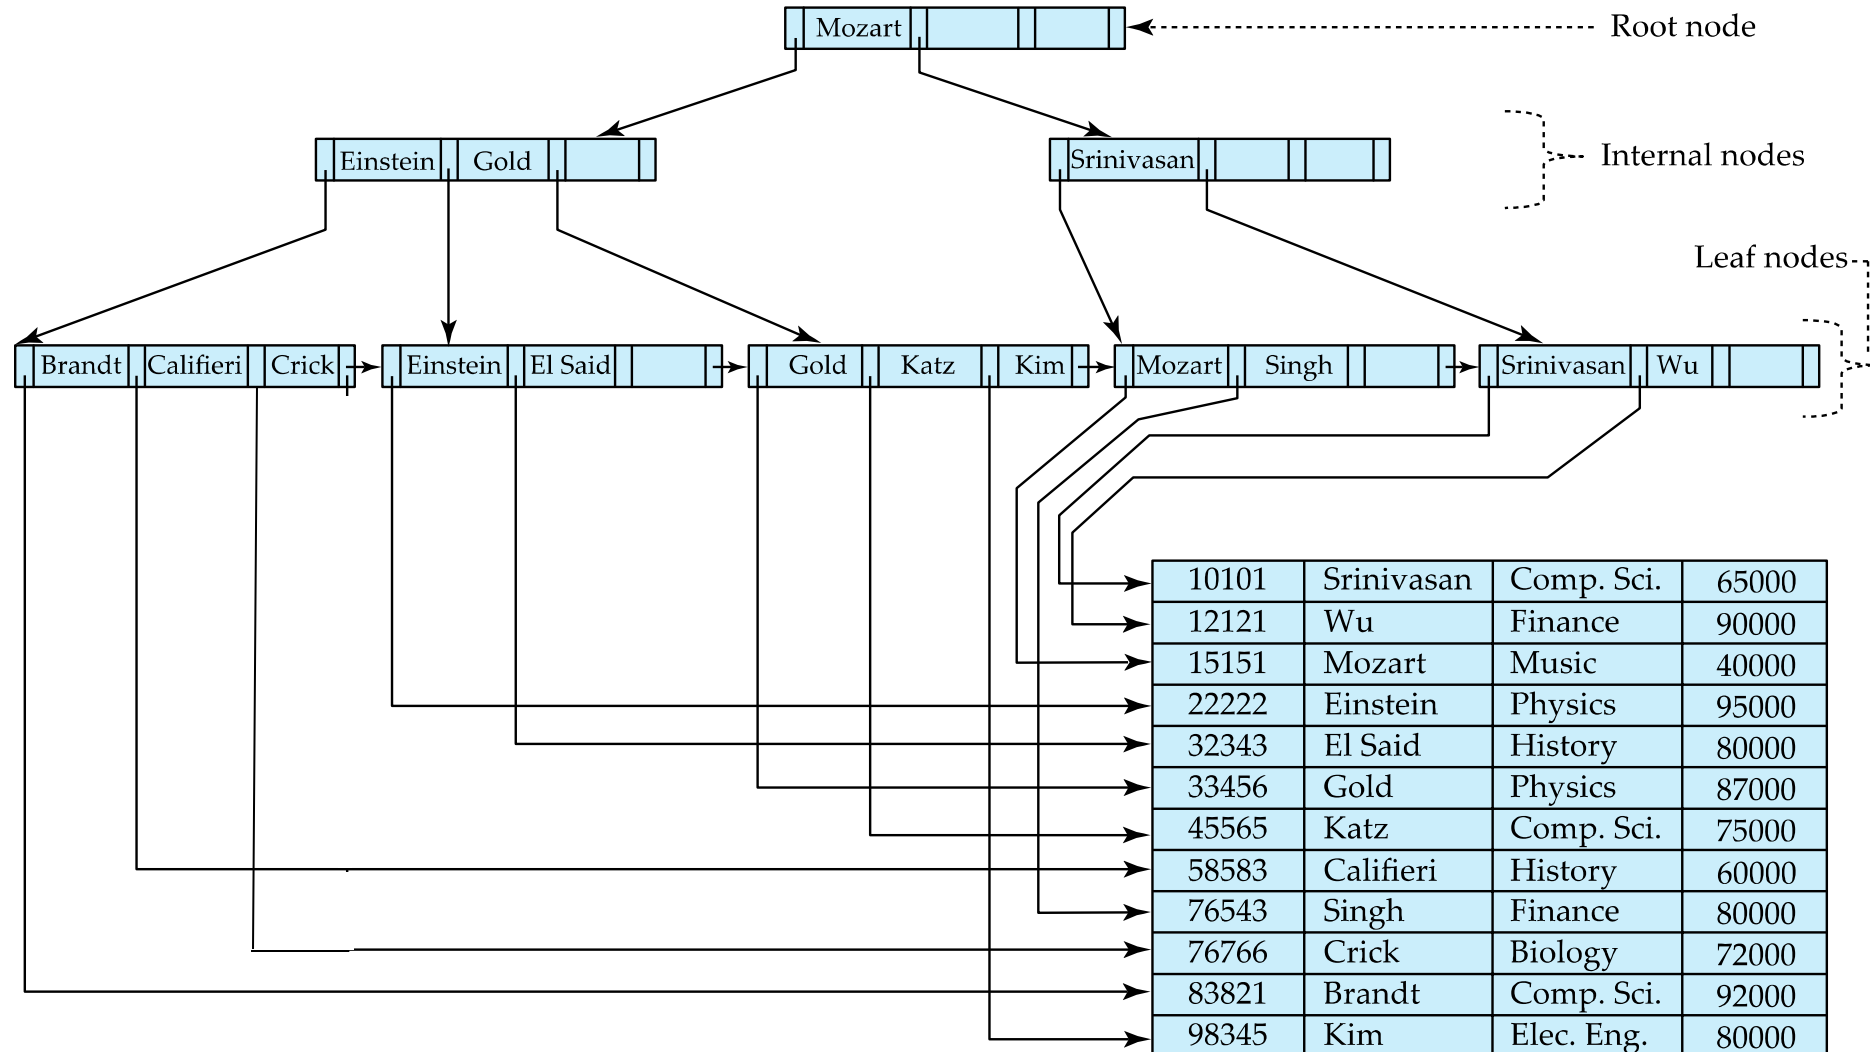


B⁺-TREE INDEX FILES


- Disadvantage of indexed-sequential files
 - Performance degrades as file grows, since many overflow blocks get created.
 - Periodic reorganization of entire file is required.
- Advantage of B⁺-tree index files:
 - Automatically reorganizes itself with small, local, changes, in the face of insertions and deletions.
 - Reorganization of entire file is not required to maintain performance.
- (Minor) disadvantage of B⁺-trees:
 - Extra insertion and deletion overhead, space overhead.
- Advantages of B⁺-trees outweigh disadvantages
 - B⁺-trees are used extensively



EXAMPLE OF B⁺-TREE



B + TREE BASIC




Source : https://youtu.be/49P_GDeMDRo

B + TREE INSERTION



B + TREE DELETION

Source : <https://youtu.be/QrbaQDSuxIM>



CREATION OF INDICES



CREATION OF INDICES

- Example
 - create index** *takes_pk* **on** *takes* (*ID*, *course_ID*, *year*, *semester*, *section*)
 - drop index** *takes_pk*
- Most database systems allow specification of type of index, and clustering.
- Indices on primary key created automatically by all databases
 - Why?
- Some database also create indices on foreign key attributes
 - Why might such an index be useful for this query:
 - *takes* ⋈ $\sigma_{name='Shankar'}$ (*student*)
- Indices can greatly speed up lookups, but impose cost on updates
 - Index tuning assistants/wizards supported on several databases to help choose indices, based on query and update workload



INDEX DEFINITION IN SQL

- Create an index:

```
create index <index-name> on <relation-name>(<attribute-list>)
```

E.g.,: **create index** *b-index* **on** *branch(branch_name)*

- Use **create unique index** to indirectly specify and enforce the condition that the search key is a candidate key is a candidate key.
 - Not really required if SQL **unique** integrity constraint is supported
- To drop an index:

```
drop index <index-name>
```
- Most database systems allow specification of type of index, and clustering.



REFERENCES

Silberschatz, Korth, and Sudarshan. *Database System Concepts* – 7th Edition. McGraw-Hill. 2019.

Slides adapted from Database System Concepts Slide.

Source: <https://www.db-book.com/db7/slides-dir/index.html>





ANY QUESTIONS?

