

Database System

09 | Concurrency Control

Tahun Ajar Ganjil 2024/2025

Oleh:
Tim Dosen

Goals of the Meeting


01

Students understand
how the working of
protocol locking




02

Students understand
how to handling
deadlock problem



OUTLINES

- Lock-Based Protocols
- Deadlock Handling



LOCK-BASED PROTOCOL



INTRODUCTION

- The fundamental properties of a transaction is **isolation**.
- When several transactions execute concurrently in the database, however, the isolation property may no longer be preserved.
- To ensure that it is, the system must control the interaction among the concurrent transactions; this control is achieved through one of a variety of mechanisms called **concurrency-control** schemes.
- There are a variety of concurrency-control schemes. No one scheme is clearly the best; each one has advantages.
- In practice, the most frequently used schemes are **two-phase locking** and **snapshot isolation**.
- In this course we will learn the **lock-based protocols**.



LOCK-BASED PROTOCOLS

- One way to ensure isolation is to require that data items be accessed in a mutually exclusive manner;
- That is, while one transaction is accessing a data item, no other transaction can modify that data item.
- The most common method used to implement this requirement is to allow a transaction to access a data item only if it is currently holding a **lock** on that item.




LOCK-BASED PROTOCOLS (CONT.)

- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks.
- Locking protocols enforce serializability by restricting the set of possible schedules.



LOCK-BASED PROTOCOLS (CONT.)

- Data items can be locked in two modes :
 1. **exclusive** (*X*) *mode*. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
 2. **shared** (*S*) *mode*. Data item can only be read. S-lock is requested using **lock-S** instruction.
- Lock requests are made to concurrency-control manager.
- Transaction can proceed only after request is granted.



LOCK-BASED PROTOCOLS (CONT.)

- **Lock-compatibility matrix**

| | S | X |
|---|-------|-------|
| S | true | false |
| X | false | false |

- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- Any number of transactions can hold shared locks on an item,
- But if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.



LOCK-BASED PROTOCOLS (CONT.)

- Example of a transaction performing locking:

T_2 : **lock-S**(A);

read (A);

unlock(A);


lock-S(B);

read (B);

unlock(B);

display(A+B)

- Locking as above is not sufficient to guarantee serializability



SCHEDULE WITH LOCK GRANTS

- This schedule is not serializable (why?)

Note:

- Grants omitted in rest of chapter. Assume grant happens just before the next instruction following lock request

| T_1 | T_2 | concurrency-control manager |
|---------------|--------------------|-----------------------------|
| lock-X(B) | | grant-X(B, T_1) |
| read(B) | | |
| $B := B - 50$ | | |
| write(B) | | |
| unlock(B) | | |
| | lock-S(A) | grant-S(A, T_2) |
| | read(A) | |
| | unlock(A) | |
| | lock-S(B) | grant-S(B, T_2) |
| | read(B) | |
| | unlock(B) | |
| | display($A + B$) | |
| lock-X(A) | | grant-X(A, T_1) |
| read(A) | | |
| $A := A + 50$ | | |
| write(A) | | |
| unlock(A) | | |



SCHEDULE WITH LOCK GRANTS (CONT.)

- Suppose that the values of accounts A and B are \$100 and \$200, respectively.
- If these two transactions are executed serially, either in the order T1, T2 or the order T2, T1, then transaction T2 will display the value \$300.
- If, however, these transactions are executed concurrently, then schedule 1, is possible.
- In this case, transaction T2 displays \$250, which is incorrect.
- The reason for this mistake is that the transaction T1 unlocked data item B too early, as a result of which T2 saw an **inconsistent state**.

| T_1 | T_2 |
|---------------|----------------|
| lock-X(B) | |
| read(B) | |
| $B := B - 50$ | |
| write(B) | |
| unlock(B) | |
| | lock-S(A) |
| | read(A) |
| | unlock(A) |
| | lock-S(B) |
| | read(B) |
| | unlock(B) |
| | display(A + B) |
| lock-X(A) | |
| read(A) | |
| $A := A + 50$ | |
| write(A) | |
| unlock(A) | |

Schedule 1



SCHEDULE WITH LOCK GRANTS (CONT.)

- Suppose now that unlocking is delayed to the end of the transaction.
Transaction T3 corresponds to T1 with unlocking delayed.
- Transaction T4 corresponds to T2 with unlocking delayed.
- T4 will not print out an inconsistent result in any of them.

T_3 : lock-X(B);
read(B);
 $B := B - 50$;
write(B);
lock-X(A);
read(A);
 $A := A + 50$;
write(A);
unlock(B);
unlock(A).

T_4 : lock-S(A);
read(A);
lock-S(B);
read(B);
display($A + B$);
unlock(A);
unlock(B).



DEADLOCK

- Consider the partial schedule

| T_3 | T_4 |
|---------------|---------------|
| lock-X(B) | |
| read(B) | |
| $B := B - 50$ | |
| write(B) | |
| | lock-S(A) |
| | read(A) |
| | lock-S(B) |
| lock-X(A) | |

- Neither T_3 nor T_4 can make progress — executing **lock-S(B)** causes T_4 to wait for T_3 to release its lock on B , while executing **lock-X(A)** causes T_3 to wait for T_4 to release its lock on A .
- Such a situation is called a **deadlock**.
 - To handle a deadlock one of T_3 or T_4 must be rolled back and its locks released.



DEADLOCK (CONT.)

- The potential for deadlock exists in most locking protocols.
- Deadlocks are a necessary evil associated with locking, if we want to avoid inconsistent states.
- Deadlocks are definitely preferable to inconsistent states, since they can be handled by rolling back transactions, whereas inconsistent states may lead to real-world problems that cannot be handled by the database system.





STARVATION

- **Starvation** is also possible if concurrency control manager is badly designed.
- For example:
 - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
 - The same transaction is repeatedly rolled back due to deadlocks.
- Concurrency control manager can be designed to prevent starvation.



THE TWO-PHASE LOCKING PROTOCOL

- A protocol which ensures conflict-serializable schedules.
- Phase 1: **Growing Phase**
 - Transaction may obtain locks
 - Transaction may not release locks
- Phase 2: **Shrinking Phase**
 - Transaction may release locks
 - Transaction may not obtain locks
- The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points** (i.e., the point where a transaction acquired its final lock).



THE TWO-PHASE LOCKING PROTOCOL (CONT.)

- Two-phase locking *does not* ensure freedom from deadlocks
- Extensions to basic two-phase locking needed to ensure recoverability of freedom from cascading roll-back
 - **Strict two-phase locking:** a transaction must hold all its exclusive locks till it commits/aborts.
 - Ensures recoverability and avoids cascading roll-backs
 - **Rigorous two-phase locking:** a transaction must hold *all* locks till commit/abort.
 - Transactions can be serialized in the order in which they commit.
- Most databases implement rigorous two-phase locking, *but refer to it as simply two-phase locking*



REFERENCES

Silberschatz, Korth, and Sudarshan. *Database System Concepts* – 7th Edition. McGraw-Hill. 2019.

Slides adapted from Database System Concepts Slide.

Source: <https://www.db-book.com/db7/slides-dir/index.html>

Elmasri, Navathe, “Fundamental of Database Systems”, Seventh Edition, Pearson, 2015.





ANY QUESTIONS?

